

Class 8 Mini-Project

Mari Williams (A15858833)

Table of contents

Data prep	1
Principal Component Analysis	3
Interpreting PCA results	5
Variance explained	8
Communicating PCA Results	9
Hierarchical clustering	9
Combining methods	13

Data prep

```
fna.data <- "WisconsinCancer.csv"  
wisc.df <- read.csv(fna.data, row.names=1)  
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	

843786	0.12780	0.17000	0.1578	0.08089		
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se	
842302	0.2419		0.07871	1.0950	0.9053	8.589
842517	0.1812		0.05667	0.5435	0.7339	3.398
84300903	0.2069		0.05999	0.7456	0.7869	4.585
84348301	0.2597		0.09744	0.4956	1.1560	3.445
84358402	0.1809		0.05883	0.7572	0.7813	5.438
843786	0.2087		0.07613	0.3345	0.8902	2.217
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	153.40	0.006399		0.04904	0.05373	0.01587
842517	74.08	0.005225		0.01308	0.01860	0.01340
84300903	94.03	0.006150		0.04006	0.03832	0.02058
84348301	27.23	0.009110		0.07458	0.05661	0.01867
84358402	94.44	0.011490		0.02461	0.05688	0.01885
843786	27.19	0.007510		0.03345	0.03672	0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302	0.03003		0.006193	25.38	17.33	
842517	0.01389		0.003532	24.99	23.41	
84300903	0.02250		0.004571	23.57	25.53	
84348301	0.05963		0.009208	14.91	26.50	
84358402	0.01756		0.005115	22.54	16.67	
843786	0.02165		0.005082	15.47	23.75	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302	184.60	2019.0		0.1622	0.6656	
842517	158.80	1956.0		0.1238	0.1866	
84300903	152.50	1709.0		0.1444	0.4245	
84348301	98.87	567.7		0.2098	0.8663	
84358402	152.20	1575.0		0.1374	0.2050	
843786	103.40	741.6		0.1791	0.5249	
	concavity_worst	concave.points_worst	symmetry_worst			
842302	0.7119		0.2654	0.4601		
842517	0.2416		0.1860	0.2750		
84300903	0.4504		0.2430	0.3613		
84348301	0.6869		0.2575	0.6638		
84358402	0.4000		0.1625	0.2364		
843786	0.5355		0.1741	0.3985		
	fractal_dimension_worst					
842302		0.11890				
842517		0.08902				
84300903		0.08758				
84348301		0.17300				
84358402		0.07678				
843786		0.12440				

```
wisc.data <- wisc.df[,-1]
diagnosis <- wisc.df$diagnosis
diagnosis <-as.factor(diagnosis)
```

Q1. How many observations are in this dataset?

```
dim(wisc.data)
```

[1] 569 30

569 observations

Q2. How many of the observations have a malignant diagnosis?

```
sum(diagnosis=="M")
```

[1] 212

212 observations are diagnosed as malignant

Q3. How many variables/features in the data are suffixed with _mean

```
length(grep("_mean", colnames(wisc.data)))
```

[1] 10

10 variables are suffixed with _mean

Principal Component Analysis

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01

```

fractal_dimension_mean           radius_se          texture_se
6.279761e-02                   4.051721e-01   1.216853e+00
perimeter_se                    area_se           smoothness_se
2.866059e+00                  4.033708e+01   7.040979e-03
compactness_se                  concavity_se      concave.points_se
2.547814e-02                   3.189372e-02   1.179614e-02
symmetry_se                     fractal_dimension_se radius_worst
2.054230e-02                   3.794904e-03   1.626919e+01
texture_worst                   perimeter_worst   area_worst
2.567722e+01                   1.072612e+02   8.805831e+02
smoothness_worst                compactness_worst concavity_worst
1.323686e-01                   2.542650e-01   2.721885e-01
concave.points_worst            symmetry_worst   fractal_dimension_worst
1.146062e-01                   2.900756e-01   8.394582e-02

```

```
apply(wisc.data, 2, sd)
```

```

radius_mean                      texture_mean      perimeter_mean
3.524049e+00                   4.301036e+00   2.429898e+01
area_mean                        smoothness_mean compactness_mean
3.519141e+02                   1.406413e-02   5.281276e-02
concavity_mean                  concave.points_mean symmetry_mean
7.971981e-02                   3.880284e-02   2.741428e-02
fractal_dimension_mean           radius_se         texture_se
7.060363e-03                   2.773127e-01   5.516484e-01
perimeter_se                    area_se           smoothness_se
2.021855e+00                   4.549101e+01   3.002518e-03
compactness_se                  concavity_se     concave.points_se
1.790818e-02                   3.018606e-02   6.170285e-03
symmetry_se                     fractal_dimension_se radius_worst
8.266372e-03                   2.646071e-03   4.833242e+00
texture_worst                   perimeter_worst   area_worst
6.146258e+00                   3.360254e+01   5.693570e+02
smoothness_worst                compactness_worst concavity_worst
2.283243e-02                   1.573365e-01   2.086243e-01
concave.points_worst            symmetry_worst   fractal_dimension_worst
6.573234e-02                   6.186747e-02   1.806127e-02

```

```
wisc.pr <- prcomp(wisc.data, center = TRUE, scale. = TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

44% of the variance is explained by PC1

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

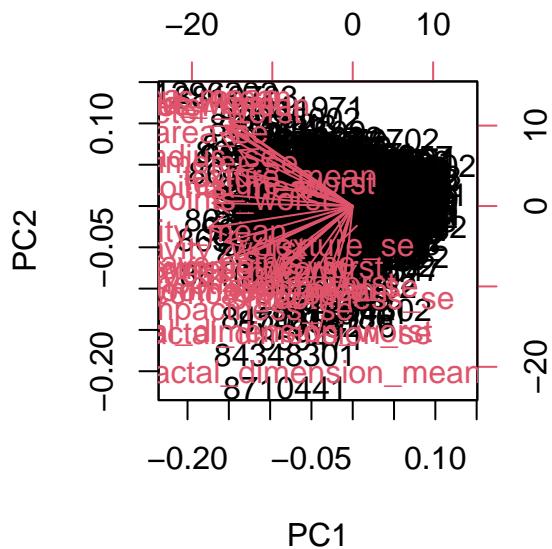
3 PCs

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PCs

Interpreting PCA results

```
biplot(wisc.pr)
```

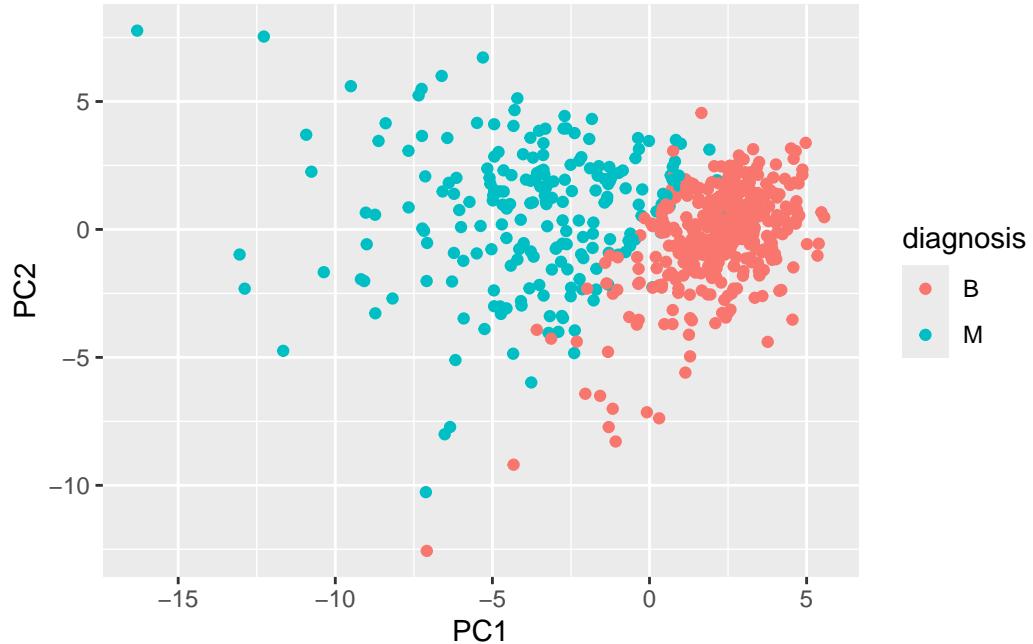


*Q7. What stands out to you about this plot? Is it easy or difficult to understand?
Why?* Text and numbers are indecipherable

```
library(ggplot2)
```

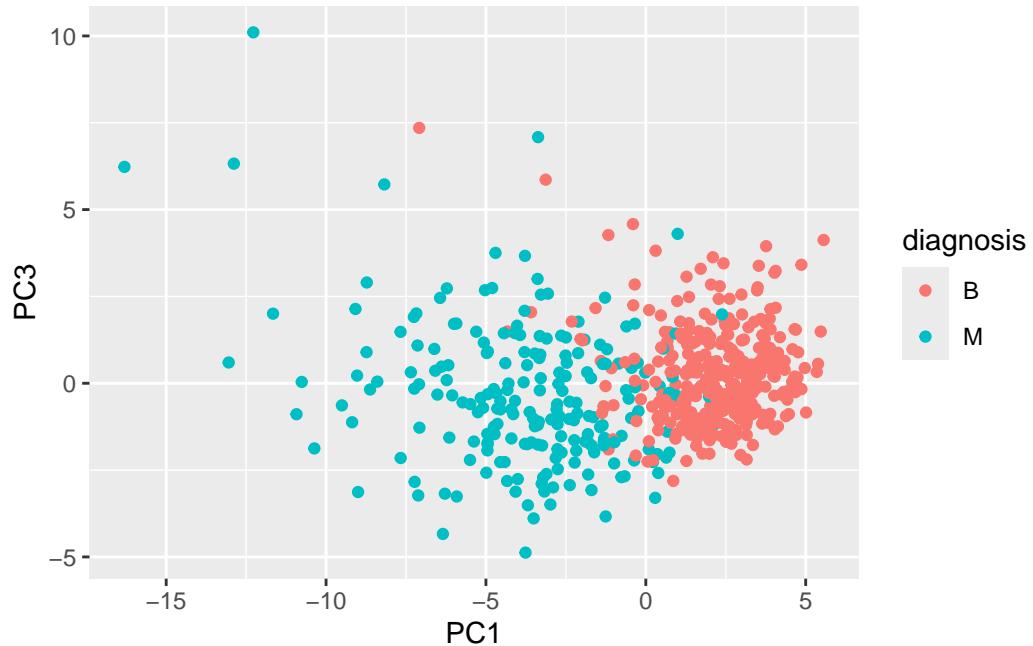
```
Warning: package 'ggplot2' was built under R version 4.3.3
```

```
ggplot(wisc.pr$x) + aes(PC1, PC2, col = diagnosis, xlab = "PC1", ylab= "PC2") + geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
ggplot(wisc.pr$x) + aes(PC1, PC3, col = diagnosis, xlab = "PC1", ylab= "PC3") + geom_point()
```



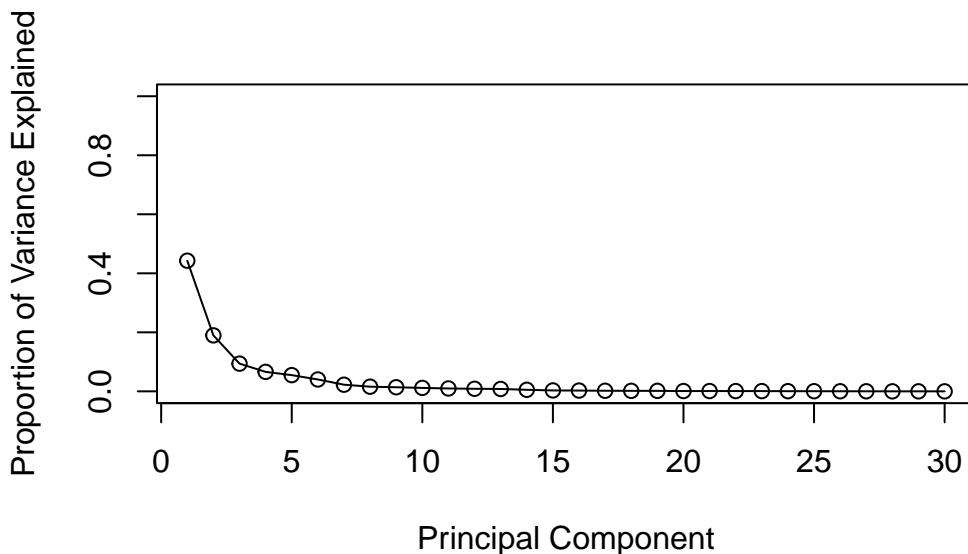
Both plots have pretty good separation between diagnoses but the first one is a little cleaner.
Similar x distributions, the y distributions are shifted.

Variance explained

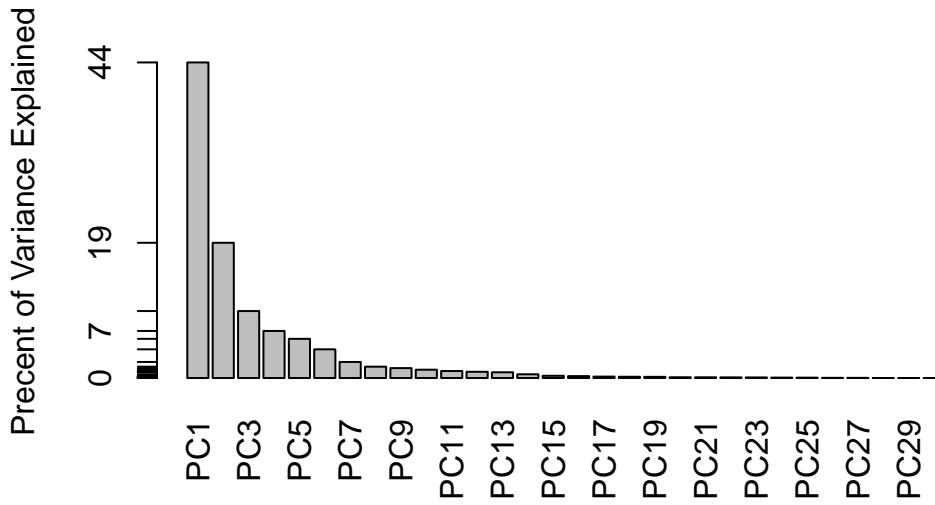
```
pr.var <- wisc.pr$sdev^2  
head(pr.var)
```

```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

```
pve <- pr.var / sum(pr.var)  
plot(pve, xlab = "Principal Component",  
     ylab = "Proportion of Variance Explained",  
     ylim = c(0, 1), type = "o")
```



```
barplot(pve, ylab = "Percent of Variance Explained",  
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)  
axis(2, at=pve, labels=round(pve,2)*100 )
```



Communicating PCA Results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC.

```
lv <- wisc.pr$rotation[,1]
lv["concave.points_mean"]
```

```
concave.points_mean
-0.2608538
```

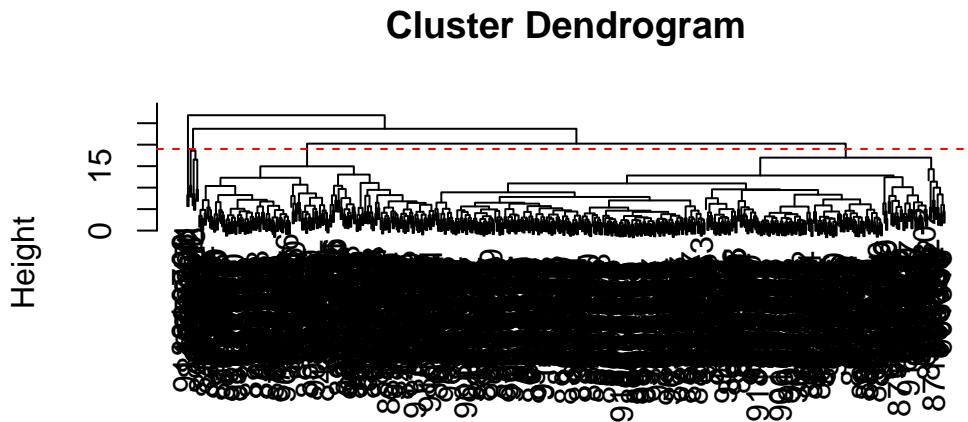
-0.261

Hierarchical clustering

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(data.dist, method= "complete")
```

Q10. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(a=19, b=0, col="red", lty=2)
```



```
data.dist
hclust (*, "complete")

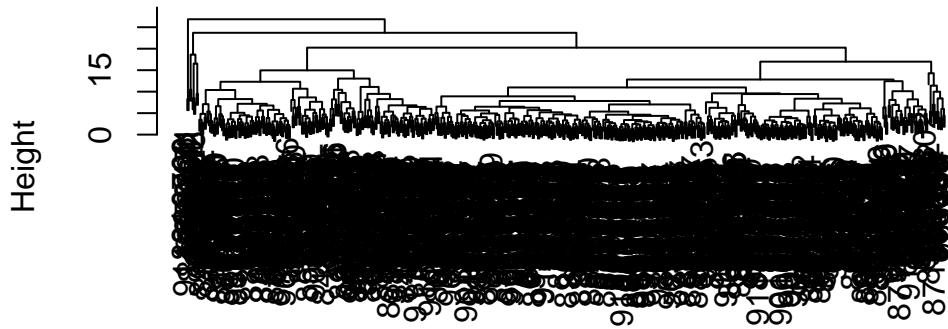
height=19
```

Q12. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

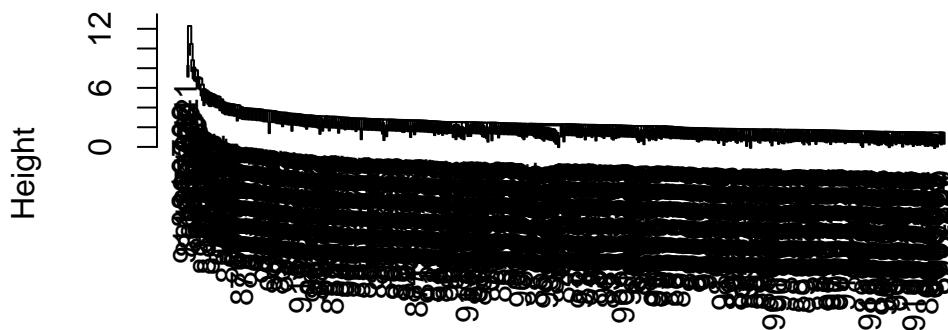
```
clust <- function(data, way) {
  wisc.hclust <- hclust(data, method= way)
  plot(wisc.hclust, xlab=way)

}
sapply(c("complete", "single", "average", "ward.D2"), function(m) clust(data.dist, m))
```

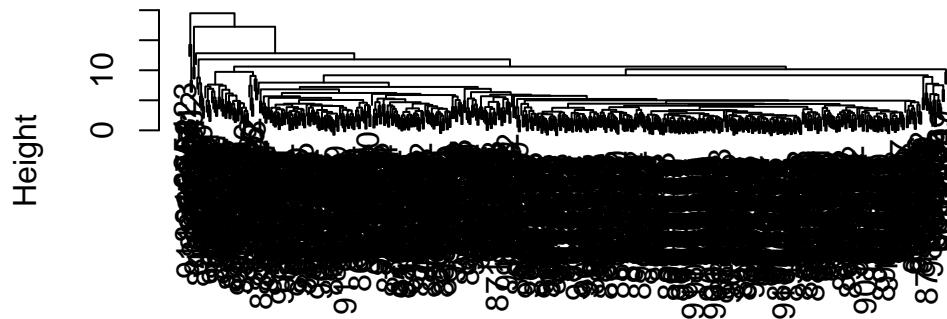
Cluster Dendrogram



Cluster Dendrogram

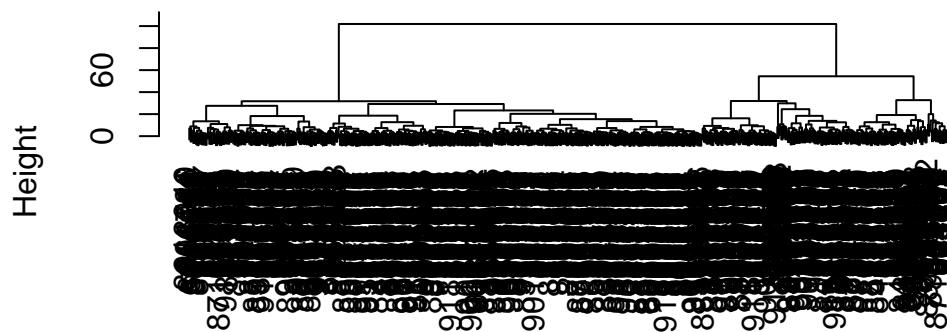


Cluster Dendrogram



average
hclust (*, "average")

Cluster Dendrogram



ward.D2
hclust (*, "ward.D2")

```
$complete  
NULL
```

```
$single  
NULL
```

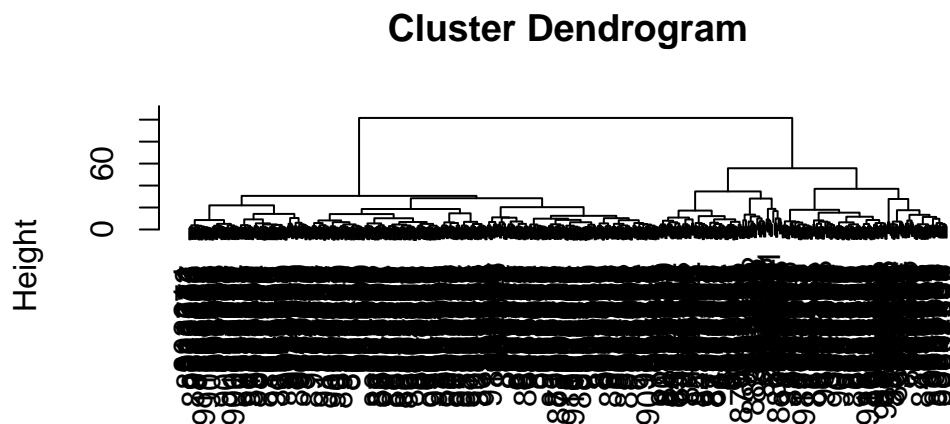
```
$average  
NULL
```

```
$ward.D2  
NULL
```

I like ward.D2 as well, because of the clean separation of the clusters

Combining methods

```
d<- dist(wisc.pr$x[,1:3])  
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7]), method= "ward.D2")  
plot(wisc.pr.hclust)
```



```
dist(wisc.pr$x[, 1:7])  
hclust (*, "ward.D2")
```

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:3]), method= "ward.D2")  
grps <- cutree(wisc.pr.hclust, k=2)  
table(grps)
```

```
grps
  1   2
203 366
```

```
table(grps, diagnosis)
```

```
diagnosis
grps   B    M
  1  24 179
  2 333 33
```

TP: 179 FP:33

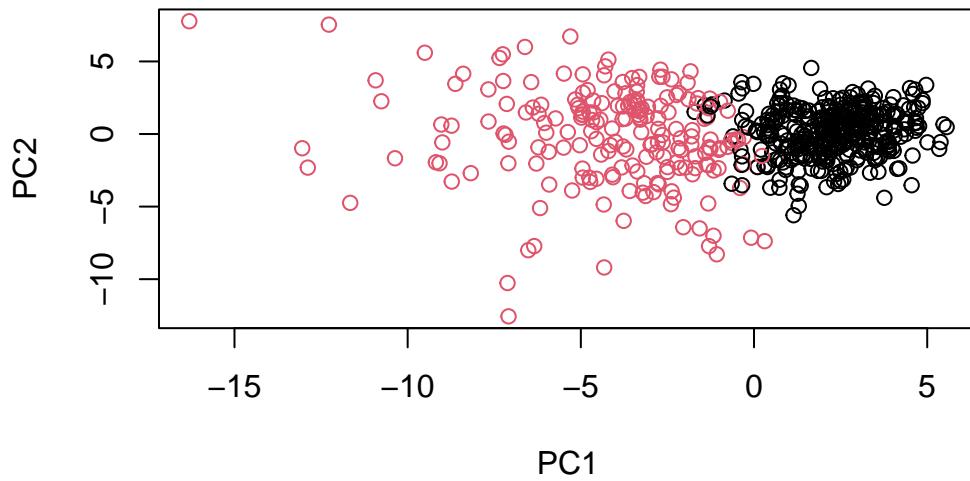
```
g <- as.factor(grps)
levels(g)
```

```
[1] "1" "2"
```

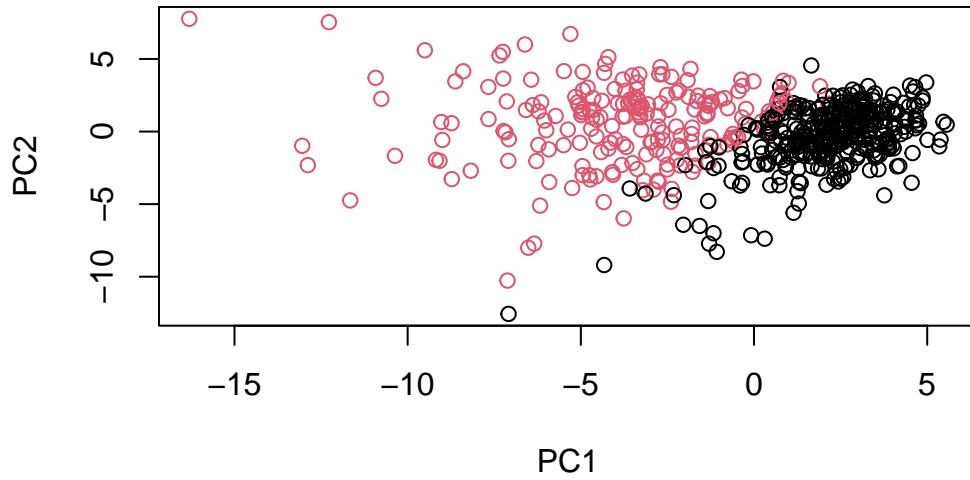
```
g <- relevel(g,2)
levels(g)
```

```
[1] "2" "1"
```

```
plot(wisc.pr$x[,1:2], col=g)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```



```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[, 1:7]), method="ward.D2")

wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
table(wisc.pr.hclust.clusters)
```

```
wisc.pr.hclust.clusters
 1   2
216 353
```

Q13. How well does the newly created model with four clusters separate out the two diagnoses?

```
table(wisc.pr.hclust.clusters, diagnosis)
```

		diagnosis	
		B	M
wisc.pr.hclust.clusters	1	28	188
	2	329	24

Pretty well! 9% of the assignments do not agree, but the rest do.

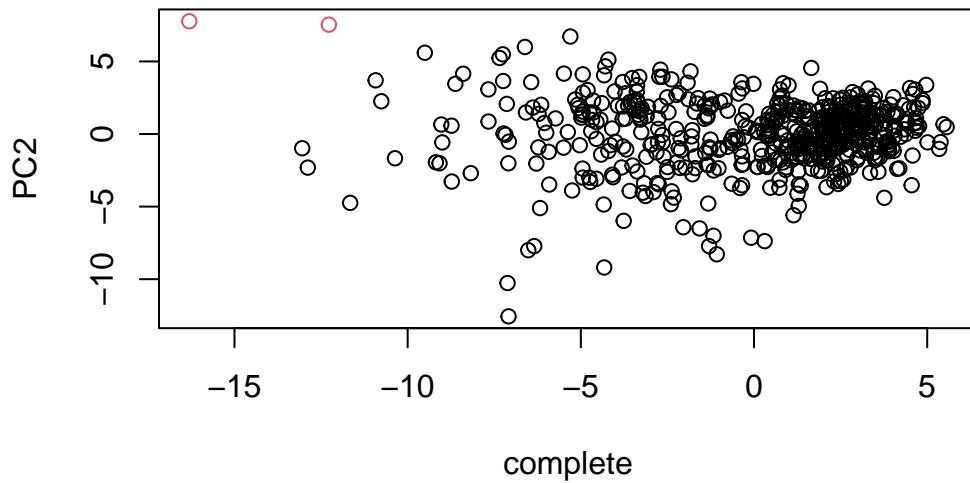
Q14. How well do the hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses.

```
clust <- function(data, way) {
  wisc.hclust <- hclust(data, method= way)
  wisc.pr.hclust.clusters <- cutree(wisc.hclust, k=2)

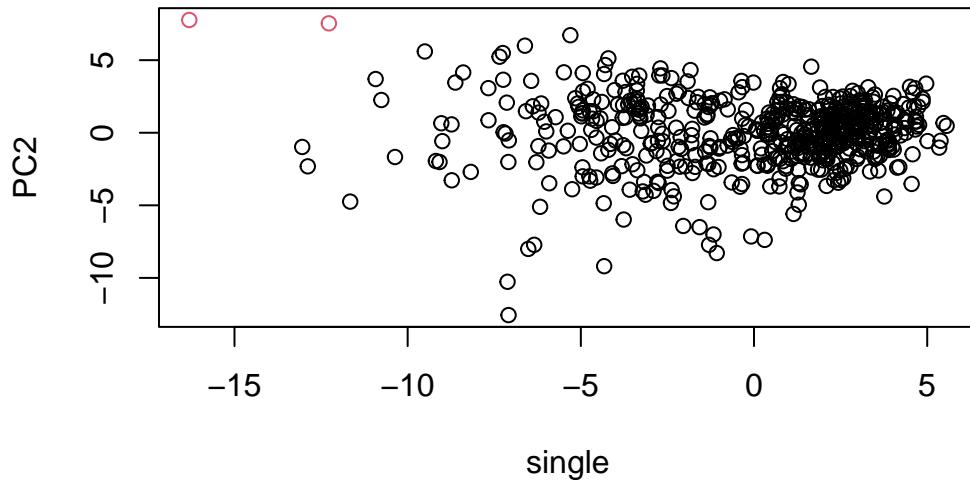
  tbl<- (table(wisc.pr.hclust.clusters, diagnosis))
  plt<- (plot(wisc.pr$x[,1:2], col=wisc.pr.hclust.clusters, xlab=way))

  return(list(table = tbl, plot = plt))
}

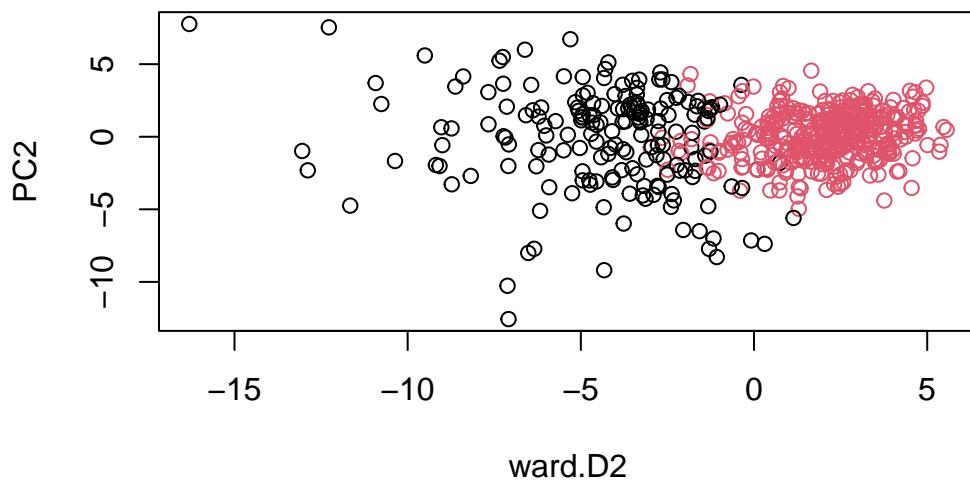
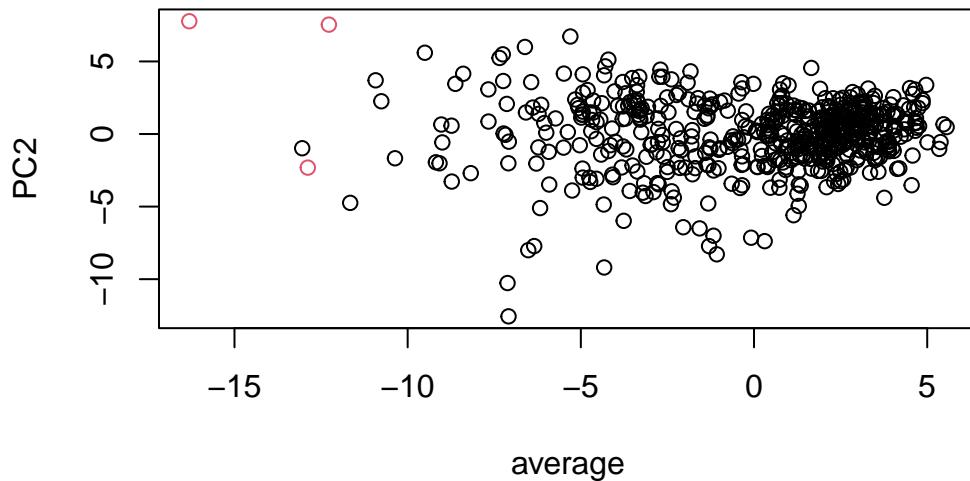
result <- sapply(c("complete", "single", "average", "ward.D2"), function(m) clust(data.dist,
```



complete



single



```
result[1,1:4]
```

```
$complete
```

```

diagnosis
wisc.pr.hclust.clusters   B   M
      1 357 210
      2   0   2

$single
diagnosis
wisc.pr.hclust.clusters   B   M
      1 357 210
      2   0   2

$average
diagnosis
wisc.pr.hclust.clusters   B   M
      1 357 209
      2   0   3

$ward.D2
diagnosis
wisc.pr.hclust.clusters   B   M
      1  20 164
      2 337  48

```

```
result$plot
```

```
NULL
```

Most of the methods except ward.D2 have very inaccurate categorization that does not overlap with the diagnosis.